

# NFSTRACE. Release notes.

Version: 0.3.0

[Introduction](#)

[Definitions](#)

[Portability](#)

[Target OS](#)

[Hardware Interfaces](#)

[Implementation Details](#)

[Command-line interface \(CLI\)](#)

[Running modes](#)

[Packets Filtration by BPF](#)

[Payload Filtration](#)

[Dump File Format](#)

[Compression](#)

[Pluggable Analysis Modules](#)

[Analysis Toolset](#)

[Operation Breakdown Analyzer \(OB\)](#)

[Performance Testing Results](#)

## Introduction

The nfstrace tool performs live Ethernet 1Gbps - 10Gbps packet's capturing and helps to determine NFS procedures in raw network traffic. It performs packet's capturing, filtration, dumping traces, compression, statistical analysis, visualization and provides API for custom analysis modules.

The nfstrace supports following protocols: Ethernet – IPv4|IPv6 – UDP|TCP – NFSv3|NFSv4.

## Definitions

SUS – Single UNIX Specification

POSIX – Portable Operating System Interface for Unix

[GPL](#) – General Public License

[Wireshark](#) – Enterprise quality tool-set for network traffic analysis.

[BPF](#) – Berkeley Packet Filter

[LSF](#) – Linux Socket Filtering, is derived from the BPF

[NIC](#) – Network Interface Card

[NFS](#) – Network File System Protocol

RPC/NFSv3/NFSv4 headers – Parts of RPC messages, which are useful in analysis of NFS operations

Payload – User's data transferred by NFS protocol. It is useless in analysis

[wsize, rsize](#) – options of NFS client connection to a NFS server

[DPI](#) – Deep Packet Inspection, nfstrace performs DPI of raw network traffic

[gnuplot](#) – CLI tool that can generate two- and three-dimensional plots of data, and data fits.

## Portability

### Target OS

The software has been developed and run as command line utility on GNU/Linux (Fedora 20, OpenSUSE 13.1, Ubuntu 13.10, Arch Linux) and FreeBSD (FreeBSD 8.4, FreeBSD 10.0). The implementation written on C++11 programming language. It uses standard POSIX interfaces and following libraries: libpthread, libpcap, libstdc++.

```
$ ldd nfstrace
linux-vdso.so.1 (0x00007ffffbf595000)
libdl.so.2 => /usr/lib/libdl.so.2 (0x00007f3088698000)
libpthread.so.0 => /usr/lib/libpthread.so.0 (0x00007f308847a000)
libpcap.so.1 => /usr/lib/libpcap.so.1 (0x00007f3088233000)
libstdc++.so.6 => /usr/lib/libstdc++.so.6 (0x00007f3087f24000)
libm.so.6 => /usr/lib/libm.so.6 (0x00007f3087c20000)
libgcc_s.so.1 => /usr/lib/libgcc_s.so.1 (0x00007f3087a0a000)
libc.so.6 => /usr/lib/libc.so.6 (0x00007f308765c000)
/lib64/ld-linux-x86-64.so.2 (0x00007f308889c000)
libusb-1.0.so.0 => /usr/lib/libusb-1.0.so.0 (0x00007f3087444000)
libnl-genl-3.so.200 => /usr/lib/libnl-genl-3.so.200 (0x00007f308723d000)
libnl-3.so.200 => /usr/lib/libnl-3.so.200 (0x00007f308701e000)
libdbus-1.so.3 => /usr/lib/libdbus-1.so.3 (0x00007f3086dd6000)
libudev.so.1 => /usr/lib/libudev.so.1 (0x00007f3088a7a000)
librt.so.1 => /usr/lib/librt.so.1 (0x00007f3086bce000)
```

The software does not use any external libraries due to portability reasons and licensing limitations. The nfstrace does not use any Wireshark libraries because they are not efficient.

### Hardware Interfaces

The nfstrace captures raw packets from an Ethernet interface via libpcap interface of Linux (LSF) or FreeBSD (BPF) implementations. Current implementation makes assumption, that libpcap delivers correct TCP and UDP packets. Assembling Ethernet frames to IP packets and IP packets defragmentation performs in the Kernel. The software have tested on direct linked workstations with 1Gbps integrated NICs (Ethernet 1000baseT/Full). The nfstrace demonstrates perfect performance even in live capturing/filtration/analysis mode.

# Implementation Details

## Command-line interface (CLI)

The nfstrace has following command-line interface:

```
$ ./nfstrace -a libbreakdown.so -h
nfstrace 0.3 (Release)
built on Linux-3.16.1-1-generic
by C++ compiler GNU 4.9.1
Usage: ./nfstrace [OPTIONS]...
  -m, --mode=drain|live|dump|stat (default:live) set runing mode
  -i, --interface=INTERFACE (default:PCAP-DEFAULT) listen interface, required for live and dump modes
  -f, --filtration=BPF (default:port 2049) a packet filtration in libpcap BPF syntax
  -s, --snaplen=0..65535 (default:65535) max length of raw captured packet. May be used ONLY FOR
UDP
  -t, --timeout=Milliseconds (default:100) set the read timeout that will be used on a capture
  -b, --bsize=MBytes (default:20) size of capturing kernel buffer
  -p, --promisc (default:true) put the interface into promiscuous mode
  -d, --direction=in|out|inout (default:inout) set the direction for which packets will be captured
  -a, --analysis=PATH#opt1,opt2=val,...(default:) specify path to analysis module and set desired
options
  -I, --ifile=PATH (default:PROGRAMNAME-BPF.pcap) input file for stat mode, the '-' means stdin
  -O, --ofile=PATH (default:PROGRAMNAME-BPF.pcap) output file for dump mode, the '-' means
stdout
  -C, --command="shell command" (default:) execute command for each dumped file
  -D, --dump-size=MBytes (default:0) size of dumping file portion, 0 = no limit
  -L, --list (default:false) list all available network interfaces
  -M, --msg-header=1..4000 (default:512) RPC message will be truncated to this limit in bytes
  -Q, --qcapacity=1..65535 (default:4096) initial queue capacity of RPC messages
  -T, --trace (default:false) print collected NFSv3 procedures, true if no modules added
  -Z, --droproot=username (default:) drops root privileges, after opening the capture device
  -v, --verbose=0|1|2 (default:1) level of print out additional information
  -h, --help (default:false) print help message and usage for modules, then exit
Usage of libbreakdown.so:
ACC - for accurate evaluation(default), MEM - for memory efficient evaluation. Options cannot be
combined
```

This example shows the output of --help (-h) operation. It shows command line interface description of the nfstrace and description of pluggable analysis modules passed by --analysis (-a) parameters.



filtration of IPv4-only traffic, use following BPF filter: 'ip and port 2049' (via -f option).

## Payload Filtration

Let see to following case: writing 1Gb file to a NFSv3 storage:

This action will be performed by some LOOKUP, READDIRPLUS, COMMIT, GETATTR, etc. and some (thousands?) WRITE procedures of NFSv3 protocol. The nfstrace must catch all of them. How many RPC/NFSv3 procedures will be perform? It depends on wsize. Wsize is an option of NFS client, it defines the number of bytes that NFS uses during the writing files to an NFS server by one WRITE procedure. When wsize is 512 Kbytes a NFS client should (in general) perform 1Gb/512Kb WRITE procedures. So, for writing 1Gb file to a NFS storage, a NFS client (with wsize=512k) should perform 2000 WRITE procedures + some other procedures like GETATTR, etc.

Each NFSv3 procedure consists by two RPC messages: Call – request from client to server and Reply – reply from server with result of requested procedure. Both RPC messages may contain data useful for analysis. Both RPC messages may contain thousands of payload bytes useless for analysis. The nfstrace must capture headers of calls and replies, and then match pairs of them to complete NFS procedures.

The '--snaplen' option set-up the number of bytes of a packet for uprising from the Kernel to user-space. In case of TCP transport layer this options is useless because TCP connection is a bidirectional stream of data (instead of UDP that is form of interchange up to 64k datagrams). In case of NFS over TCP, the nfstrace must capture whole packets and copy them to user-space from the Kernel for DPI and performing NFS statistics analysis.

Finally, the nfstrace must filtrate whole NFS traffic passed from the Kernel to user-space and detects RPC/NFS message headers (up to 4Kbytes) within gigabytes of network traffic.

Detected headers are copied into internal buffers (or dumped to a .pcap file) for statistics analysis.

The key principle of Filtration in the nfstrace is “**Discard payload ASAP**”.

The Filtration module works in separate thread that captures packets from network interface by libpcap. It matches a packet to a related session (TCP or UDP) and performs reassembling TCP flow from TCP segment of a packet. After that, a part of packet will be pass to RPCFiltrator. In case of NFSv4, full packet will be pass to RPCFiltrator, because it consists of several NFSv4 operations.

There are two RPCFiltrator in one TCP session. Both of them known state of current RPC

message in related TCP flow. They can detect RPC messages and perform action on a packet – discard it or collect it for analysis.

The ‘-b’ option of nfstrace set-up size of capturing buffer in the Kernel of operating system in megabytes. This option is very critical for capturing performance.

The wsize and rsize of NFS connections are meaningful for filtration and analysis performance too.

## Dump File Format

The nfstrace uses [libpcap file format](#) for input and output files. Filtered packets are saved in .pcap format too. Any external tool (like Wireshark) may be used to inspect filtered traces. Filtrated headers of RPC messages are passed from Filtration to Analysis modules in raw data buffers via Queue of Analysis module.

## Compression

The compression and decompression performs by external tools. It is very similar to tcpdump -z option. For example:

```
$ sudo ./nfstrace --mode=dump -i eth0 -f "ip and port 2049" -O dump.pcap --dump-size=10 -C "bzip2 -f -9"
Capture from interface: eth0
  BPF filter   : ip and port 2049
  snapshot len: 65535 bytes
  read timeout: 100 ms
  buffer size  : 20971520 bytes
  promiscuous mode: on
  capture traffic : inout
Dump packets to file: dump.pcap
  file rotation size: 10485760 bytes
  file rotation command: [bzip2 -f -9]
Note: It's potentially unsafe to run this program as root without dropping root privileges.
Note: Use -Z=username option for dropping root privileges when you run this program as user with root
privileges.
Processing packets. Press CTRL-C to quit and view results.
```

This command run nfstrace in dumping mode (--mode=dump). Capturing from network interface requires super-user privileges, so – sudo. In this mode, the application performs live capturing packets from network interface eth0 (-i eth0) and filtration NFS procedures (calls + replies). Then it dumps filtered packets that contains RPC/NFS headers (or parts of these headers). Filtration in the Kernel performs by BPF filter (-f "ip and port 2049"). It means that we are interested in any TCP or UDP packets send to or from port 2049(default port for NFS servers) over IPv4. This command specifies the output file dump.pcap (-O dump.pcap). In case of the output file will be huge (tons of Mb) it will be divided to parts by 10Mb (--dump-size=10).

When the part's dumping is compete, the dumping will be continued to another file (dump.pcap-1, dump.pcap-2, ... dump.pcap-N) and nfstrace spawns child process (by fork() & exec() calls) that executes command (passed by -C "bzip2 -f -9") "bzip2 -f -9 dump.pcap" over

each dumped part.

After dumping interruption via Control-C(SIGINT) the application closes with flushing captured data and we get a set of compressed parts of filtered traffic:

```
$ ls -la
total 1780
drwxrwxr-x 6 nst nst 4096 Apr 3 13:46 .
drwxrwxr-x 10 nst nst 4096 Mar 27 13:49 ..
-rw-r--r-- 1 root root 79711 Apr 3 13:46 dump.pcap-10.bz2
-rw-r--r-- 1 root root 128455 Apr 3 13:45 dump.pcap-1.bz2
-rw-r--r-- 1 root root 129402 Apr 3 13:45 dump.pcap-2.bz2
-rw-r--r-- 1 root root 132811 Apr 3 13:45 dump.pcap-3.bz2
-rw-r--r-- 1 root root 127140 Apr 3 13:45 dump.pcap-4.bz2
-rw-r--r-- 1 root root 130819 Apr 3 13:45 dump.pcap-5.bz2
-rw-r--r-- 1 root root 135586 Apr 3 13:45 dump.pcap-6.bz2
-rw-r--r-- 1 root root 130894 Apr 3 13:45 dump.pcap-7.bz2
-rw-r--r-- 1 root root 129277 Apr 3 13:45 dump.pcap-8.bz2
-rw-r--r-- 1 root root 129772 Apr 3 13:46 dump.pcap-9.bz2
-rw-r--r-- 1 root root 127589 Apr 3 13:45 dump.pcap.bz2
```

These compressed parts have filtered data and may be join to one .pcap file by:

```
$ ls dump.pcap*.bz2 | sort -n -t - -k 2 | xargs bzipcat > dump.pcap
```

The ls and sort commands sort file names of parts right order and a pipe passes them to bzipcat tool to decompressing parts to one huge file dump.pcap.

dump.pcap file may be open in any external tool that know .pcap format, f.e. – Wireshark. Only the dump.pcap.bz2 has [.pcap file header](#), other parts have only data and they are not be recognized as .pcap file by external tool.

The compressed parts of dump can be decompressed by bzipcat tool to stdin of nfsrace for offline analysis by an analyzer module or prints-out to console:

```
$ ls dump.pcap*.bz2 | sort -n -t - -k 2 | xargs bzipcat | ./nfstrace --mode=stat -I - -T
```

This command runs bzipcat tool, which decompress parts, merges them to one stream and put in to a pipe. Nfstrace runs in statistic mode (--mode=stat) for reading data from that pipe via stdin (-I -) and prints (by -T) all detected NFS procedures (call+reply) to console:

```
...
10.6.137.109:876 --> 10.6.137.113:2049 [UDP] WRITE
    CALL [ file: 10007010...9335c1f4 offset: 2024476672 count: 32768 stable: UNSTABLE ]
    REPLY [ status: OK ]
10.6.137.109:876 --> 10.6.137.113:2049 [UDP] COMMIT
    CALL [ file: 10007010...9335c1f4 offset: 0 count: 0 ]
    REPLY [ status: OK ]
...
```

The --verbose=2 option enables whole output:

```
...
10.6.137.109:876 --> 10.6.137.113:2049 [UDP] XID: 3875993544 RPC version: 2 RPC program: 100003 version:
```

3 WRITE

CALL [ file: 10007010ad108000000000002b960d115ac704d6c92360d50abcf7f8921080009335c1f4 offset: 2024476672 count: 32768 stable: UNSTABLE data: skipped on filtration ]

REPLY [ status: OK file\_wcc: before: void after: attributes: type: REG mode: OWNER\_READ  
OWNER\_WRITE GROUP\_READ GROUP\_WRITE OTHER\_READ nlink: 1 uid: 1000 gid: 1000 size: 2097152000 used:  
2097156096 rdev: specdata1: 0 specdata2: 0 fsid: 16302950409496320814 fileid: 524585 atime: seconds:  
1395737557 nseconds: 372964699 mtime: seconds: 1396521963 nseconds: 558634635 ctime: seconds: 1396521963  
nseconds: 558634635 count: 32768 committed: UNSTABLE verf: 53314027000b8fb4 ]  
10.6.137.109:876 --> 10.6.137.113:2049 [UDP] XID: 3892770760 RPC version: 2 RPC program: 100003 version:  
3 COMMIT

CALL [ file: 10007010ad108000000000002b960d115ac704d6c92360d50abcf7f8921080009335c1f4 offset: 0  
count: 0 ]

REPLY [ status: OK file\_wcc: before: void after: attributes: type: REG mode: OWNER\_READ  
OWNER\_WRITE GROUP\_READ GROUP\_WRITE OTHER\_READ nlink: 1 uid: 1000 gid: 1000 size: 2097152000 used:  
2097156096 rdev: specdata1: 0 specdata2: 0 fsid: 16302950409496320814 fileid: 524585 atime: seconds:  
1395737557 nseconds: 372964699 mtime: seconds: 1396521963 nseconds: 558634635 ctime: seconds: 1396521963  
nseconds: 558634635 verf: 53314027000b8fb4 ]

...



## Pluggable Analysis Modules

The API of nfstrace provides C++ headers with definition of IAnalyzer interface and couple of related types and NFS data structures and functions. An IAnalyzer interface is a set of NFS handlers that will be called by Analysis module for each NFS procedure. All definitions of types and constants, required for development analysis modules, are defined in files in nfstrace/api/ directory. They should be included by <nfstrace/api/plugin\_api.h> header.

A pluggable analysis module should be a dynamically linked shared object and should export following C functions:

```
const char* usage (); // return description of expected opts for create(opts)
IAnalyzer*  create (const char* opts); // create and return an instance of an Analyzer
void        destroy (IAnalyzer* instance); // destroy created instance of an Analyzer
```

and should contain the following macro after declaration of these functions for proper usage:

```
NST_PLUGIN_ENTRY_POINTS (&usage, &create, &destroy)
```

The usage() function must return C-string with description of module and required parameters for creation an instance of analyzer, this string will be shown in output of "--help" option.

The IAnalyzer\* create(const char\* opts) must create and return an instance of analyzer accordingly to passed options.

The void destroy(IAnalyzer\* instance) must destroy previously created analyzer instance and performs required clean-up (for instance - close connection to a database).

All existing analyzers are implemented as pluggable analysis modules and may be attached to nfstrace via -a option.

# Analysis Toolset

## Operation Breakdown Analyzer (OB)

The OB implemented as pluggable analysis module libbreakdown.so.

The NFS procedure latency have calculate as difference between timestamp of come-in last TCP or UDP packet of header of RPC/NFS reply-message and timestamp of come-in last TCP or UDP packet of header of related NFS call-message.

[Standard Deviation](#) of latency may be calculate by [two algorithms](#). Two-pass algorithm generates correct standard deviation value but requires a lot of memory for storage all latencies till final computation. One-pass algorithm is memory-efficient but it can accumulate computation error in case of large number of small latencies. OB implements both algorithms. They may be chosen by passing parameter while attaching OB analyzer to nfstrace.

```
$ ./nfstrace -a libbreakdown.so -h
nfstrace 0.3.0 (Release)
built on Linux-3.16.1-1-generic
by C++ compiler GNU 4.9.1
Usage: ./nfstrace [OPTIONS]...
...
Usage of libbreakdown.so:
ACC - for accurate evaluation(default), MEM - for memory efficient evaluation. Options cannot be
combined
```

MEM means [Knuth's one-pass algorithm](#). ACC means stable two-pass algorithm. ACC is default.

Let see example of OB usage. bzcata tool extracts collected trace from compressed archive to pipe. This archive contains traces of packets of NFSv4 sessions over TCP protocol.

```
$ bzcata ../traces/eth-ipv4-tcp-nfsv4.pcap.bz2 | ./nfstrace -m stat -I - -a analyzers/libbreakdown.so#ACC
Log folder: /tmp/nfstrace
Loading module: 'analyzers/libbreakdown.so' with args: [ACC]
Read packets from: -
  datalink: EN10MB (Ethernet)
  version: 2.4
Note: It's potentially unsafe to run this program as root without dropping root privileges.
Note: Use -Z=username option for dropping root privileges when you run this program as user with root
privileges.
Processing packets. Press CTRL-C to quit and view results.
Detect session 10.6.137.47:903 --> 10.6.137.113:2049 [TCP]
### Breakdown analyzer ###
NFSv3 total procedures: 0. Per procedure:
NULL          0      0%
GETATTR       0      0%
SETATTR       0      0%
LOOKUP        0      0%
```

ACCESS	0	0%
READLINK	0	0%
READ	0	0%
WRITE	0	0%
CREATE	0	0%
MKDIR	0	0%
SYMLINK	0	0%
MKNOD	0	0%
REMOVE	0	0%
RMDIR	0	0%
RENAME	0	0%
LINK	0	0%
READDIR	0	0%
READDIRPLUS	0	0%
FSSTAT	0	0%
FSINFO	0	0%
PATHCONF	0	0%
COMMIT	0	0%

NFSv4 total procedures: 1607. Per procedure:

NULL	0	0.00%
COMPOUND	1607	100.00%

NFS4 total operations: 4819. Per operation:

ILLEGAL	0	0.00%
ACCESS	1	0.02%
CLOSE	1	0.02%
COMMIT	1	0.02%
CREATE	0	0.00%
DELEGPURGE	0	0.00%
DELEGRETURN	0	0.00%
GETATTR	1604	33.28%
GETFH	1	0.02%
LINK	0	0.00%
LOCK	0	0.00%
LOCKT	0	0.00%
LOCKU	0	0.00%
LOOKUP	1	0.02%
LOOKUPP	0	0.00%
NVERIFY	0	0.00%
OPEN	1	0.02%
OPENATTR	0	0.00%
OPEN_CONFIRM	1	0.02%
OPEN_DOWNGRADE	0	0.00%
PUTFH	1607	33.35%
PUTPUBFH	0	0.00%
PUTROOTFH	0	0.00%
READ	0	0.00%
READDIR	0	0.00%
READLINK	0	0.00%
REMOVE	0	0.00%
RENAME	0	0.00%
RENEW	0	0.00%
RESTOREFH	0	0.00%
SAVEFH	0	0.00%
SECINFO	0	0.00%
SETATTR	1	0.02%
SETCLIENTID	0	0.00%

```

SETCLIENTID_CONFIRM      0   0.00%
VERIFY                   0   0.00%
WRITE                   1600 33.20%
RELEASE_LOCKOWNER        0   0.00%
GET_DIR_DELEGATION       0   0.00%

```

Per connection info:

Session: 10.6.137.47:903 --> 10.6.137.113:2049 [TCP]

Total procedures: 1607. Per procedure:

```

NULL          Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
COMPOUND      Count: 1607 (100.00%) Min: 0.000 Max: 0.990 Avg: 0.006 StDev: 0.02459829

```

Total operations: 4819. Per operation:

```

ILLEGAL      Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
ACCESS       Count: 1 ( 0.02%) Min: 0.001 Max: 0.001 Avg: 0.001 StDev: 0.00000000
CLOSE        Count: 1 ( 0.02%) Min: 0.001 Max: 0.001 Avg: 0.001 StDev: 0.00000000
COMMIT       Count: 1 ( 0.02%) Min: 0.990 Max: 0.990 Avg: 0.990 StDev: 0.00000000
CREATE       Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
DELEGPURGE   Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
DELEGRETURN  Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
GETATTR      Count: 1604 ( 33.28%) Min: 0.001 Max: 0.018 Avg: 0.005 StDev: 0.00080875
GETFH        Count: 1 ( 0.02%) Min: 0.018 Max: 0.018 Avg: 0.018 StDev: 0.00000000
LINK         Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
LOCK         Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
LOCKT        Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
LOCKU        Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
LOOKUP       Count: 1 ( 0.02%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
LOOKUPP      Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
NVERIFY      Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
OPEN         Count: 1 ( 0.02%) Min: 0.018 Max: 0.018 Avg: 0.018 StDev: 0.00000000
OPENATTR     Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
OPEN_CONFIRM Count: 1 ( 0.02%) Min: 0.001 Max: 0.001 Avg: 0.001 StDev: 0.00000000
OPEN_DOWNGRADE Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
PUTFH        Count: 1607 ( 33.35%) Min: 0.000 Max: 0.990 Avg: 0.006 StDev: 0.02459829
PUTPUBFH     Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
PUTROOTFH    Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
READ         Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
READDIR      Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
READLINK    Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
REMOVE       Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
RENAME       Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
RENEW        Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
RESTOREFH    Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
SAVEFH       Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
SECINFO      Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
SETATTR      Count: 1 ( 0.02%) Min: 0.016 Max: 0.016 Avg: 0.016 StDev: 0.00000000
SETCLIENTID  Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
SETCLIENTID_CONFIRM Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
VERIFY       Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
WRITE        Count: 1600 ( 33.20%) Min: 0.002 Max: 0.007 Avg: 0.005 StDev: 0.00067742
RELEASE_LOCKOWNER Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
GET_DIR_DELEGATION Count: 0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000

```

Filtration is done

The OB produces .dat files for each detected NFS session:


```
$ ls -la
total 109632
drwxrwxr-x 6 nst nst      4096 Apr  3 14:08 .
drwxrwxr-x 10 nst nst      4096 Mar 27 13:49 ..
drwxrwxr-x 3 nst nst      4096 Apr  2 13:25 analyzers
-rw-rw-r-- 1 nst nst      1.3K Sep 24 08:41 breakdown_10.6.137.47:903 --> 10.6.137.113:2049 [TCP].dat
```

OB has three optional modules for demonstration visualization via gnuplot:

- breakdown\_nfsv3.plt – gnuplot script for a histogram generation (NFSv3);
- breakdown\_nfsv4.plt – gnuplot script for a histogram generation (NFSv4);
- nst.sh – “glue” for passing a .dat file to breakdown.plt and invocation gnuplot tool.

The example of invocation:

```
$ ../analyzers/nst.sh -a ../analyzers/breakdown_nfsv4.plt -d . -p 'breakdown_10.6.137.47:903*.dat' -v
gnuplot -e "i_file='./breakdown_10.6.137.47:903 --> 10.6.137.113:2049
[TCP].dat';o_file='./breakdown_10.6.137.47:903 --> 10.6.137.113:2049 [TCP].dat.png'"
../analyzers/breakdown_nfsv4.plt
```



Picture 2 - breakdown\_10.6.137.47:903 --> 10.6.137.24:2049 [TCP].dat.png

## Filtration Performance Testing Results

The `nfstrace` tool can be used to perform live traffic capturing, NFS procedures filtration and statistics analysis by pluggable analysis modules. It can be ran by following command:

```
$ sudo ./nfstrace --mode=live --interface=eth0 -b 80 -a libbreakdown.so
```

The live traffic capturing from network interface requires super-user privileges, so use `sudo`. The `nfstrace` runs in line mode (`--mode=live`) to capture packets from network interface `eth0` (`--interface=eth0`) and sets capturing buffer in the Kernel to 80Mb (`-b 80`). Finally analyzer is attached (`-a libbreakdown.so`).

Test-environment is a two development workstations with 1G Ethernet NICs connected by 1G network switch. Client workstation mounts NFS shared folder with `wsiz=32k` on TCP protocol.

```
$ mount
epbym1nw1962t6:/home/nst/nfs-share on /mnt/nfs-ubuntu13 type nfs
(rw,relatime,vers=3,rsize=32768,wsiz=32768,namlen=255,hard,proto=tcp,timeo=600,retrans=2,sec=sys,mounta
ddr=10.6.137.113,mountvers=3,mountport=42985,mountproto=tcp,local_lock=none,addr=10.6.137.113)
```

Machine 2-client (10.6.137.109) will write 16Gb data to Machine 1-server (10.6.137.113) via `dd` command (`dd if=/dev/zero of=/mnt/nfs-ubuntu13/zeros bs=1M count=16000`) to a NFS shared folder.

Machine 2-client generates NFS traffic to Machine 1-server.

Machine 1-server runs `nfstrace` and `nfsstat -Z` tool for verification. The log of it has written to `nfstrace.log` file.

Due to `wsiz=32Kb` and invocation of `dd if=/dev/zero of=/mnt/nfs-ubuntu13/zeros bs=1M count=16000` we expect 512000 WRITE procedures.

```
> dd if=/dev/zero of=/mnt/nfs-ubuntu13/zeros bs=1M count=16000
16000+0 records in
16000+0 records out
16777216000 bytes (17 GB) copied, 218.17 s, 76.9 MB/s
```

Results from `nfsstat -Z` on the NFS server:

```
$ nfsstat -Z
Collecting statistics; press CTRL-C to view results from interval (i.e., from pause to CTRL-C).
^CSignal 2 received; displaying (only) statistics gathered over the last 4 minutes, 8 seconds:
```

Server rpc stats:

calls	badcalls	badclnt	badauth	xdrcll
512962	0	0	0	0

Server nfs v3:

null	getattr	setattr	lookup	access	readlink
0	0% 1	0% 1	0% 0	0% 2	0% 0 0%
read	write	create	mkdir	symlink	mknod
0	0% 512913	99% 0	0% 0	0% 0	0% 0 0%
remove	rmdir	rename	link	readdir	readdirplus

0	0% 0	0% 0	0% 0	0% 0	0% 0	0%
fsstat	fsinfo	pathconf	commit			
0	0% 0	0% 0	0% 45	0%		

After exit by Control-C from nfstrace we have following results:

```
$ sudo ./nfstrace --mode=live --interface=eth0 -b 80 -a libbreakdown.so
Log folder: /tmp/nfstrace
Loading module: 'libbreakdown.so' with args: []
Capture from interface: eth0
  BPF filter   : ip and port 2049
  snapshot len: 65535 bytes
  read timeout: 100 ms
  buffer size  : 83886080 bytes
  promiscuous  mode: on
  capture traffic : inout
Note: It's potentially unsafe to run this program as root without dropping root privileges.
Note: Use -Z=username option for dropping root privileges when you run this program as user with root
privileges.
Processing packets. Press CTRL-C to quit and view results.
Detect session 10.6.137.109:743 --> 10.6.137.113:2049 [TCP]
^C### Breakdown analyzer ###
NFSv3 total procedures: 512962. Per procedure:
NULL          0  0.00%
GETATTR       1  0.00%
SETATTR       1  0.00%
LOOKUP        0  0.00%
ACCESS        2  0.00%
READLINK     0  0.00%
READ          0  0.00%
WRITE      512913 99.99%
CREATE        0  0.00%
MKDIR         0  0.00%
SYMLINK       0  0.00%
MKNOD        0  0.00%
REMOVE       0  0.00%
RMDIR        0  0.00%
RENAME        0  0.00%
LINK         0  0.00%
REaddir       0  0.00%
REaddirPLUS  0  0.00%
FSSTAT       0  0.00%
FSINFO       0  0.00%
PATHCONF     0  0.00%
COMMIT       45  0.01%
Per connection info:
Session: 10.6.137.109:743 --> 10.6.137.113:2049 [TCP]
Total procedures: 512962. Per procedure:
NULL          Count:  0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
GETATTR       Count:  1 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
SETATTR       Count:  1 ( 0.00%) Min: 0.441 Max: 0.441 Avg: 0.441 StDev: 0.00000000
LOOKUP        Count:  0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
ACCESS        Count:  2 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00003818
READLINK     Count:  0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
READ          Count:  0 ( 0.00%) Min: 0.000 Max: 0.000 Avg: 0.000 StDev: 0.00000000
WRITE         Count: 512913 ( 99.99%) Min: 0.000 Max: 0.924 Avg: 0.004 StDev: 0.04471008
```

CREATE	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
MKDIR	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
SYMLINK	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
MKNOD	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
REMOVE	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
RMDIR	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
RENAME	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
LINK	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
READDIR	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
READDIRPLUS	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
FSSTAT	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
FSINFO	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
PATHCONF	Count:	0 ( 0.00%)	Min: 0.000	Max: 0.000	Avg: 0.000	StDev: 0.00000000
COMMIT	Count:	45 ( 0.01%)	Min: 0.000	Max: 2.765	Avg: 1.271	StDev: 1.21419871

NFSv4 total procedures: 0. Per procedure:

NULL	0	0%
COMPOUND	0	0%

NFS4 total operations: 0. Per operation:

ILLEGAL	0	0%
ACCESS	0	0%
CLOSE	0	0%
COMMIT	0	0%
CREATE	0	0%
DELEGPURGE	0	0%
DELEGRETURN	0	0%
GETATTR	0	0%
GETFH	0	0%
LINK	0	0%
LOCK	0	0%
LOCKT	0	0%
LOCKU	0	0%
LOOKUP	0	0%
LOOKUPP	0	0%
NVERIFY	0	0%
OPEN	0	0%
OPENATTR	0	0%
OPEN_CONFIRM	0	0%
OPEN_DOWNGRADE	0	0%
PUTFH	0	0%
PUTPUBFH	0	0%
PUTROOTFH	0	0%
READ	0	0%
READDIR	0	0%
READLINK	0	0%
REMOVE	0	0%
RENAME	0	0%
RENEW	0	0%
RESTOREFH	0	0%
SAVEFH	0	0%
SECINFO	0	0%
SETATTR	0	0%
SETCLIENTID	0	0%
SETCLIENTID_CONFIRM	0	0%
VERIFY	0	0%
WRITE	0	0%
RELEASE_LOCKOWNER	0	0%



GET\_DIR\_DELEGATION 0 0%

Statistic from interface: eth0

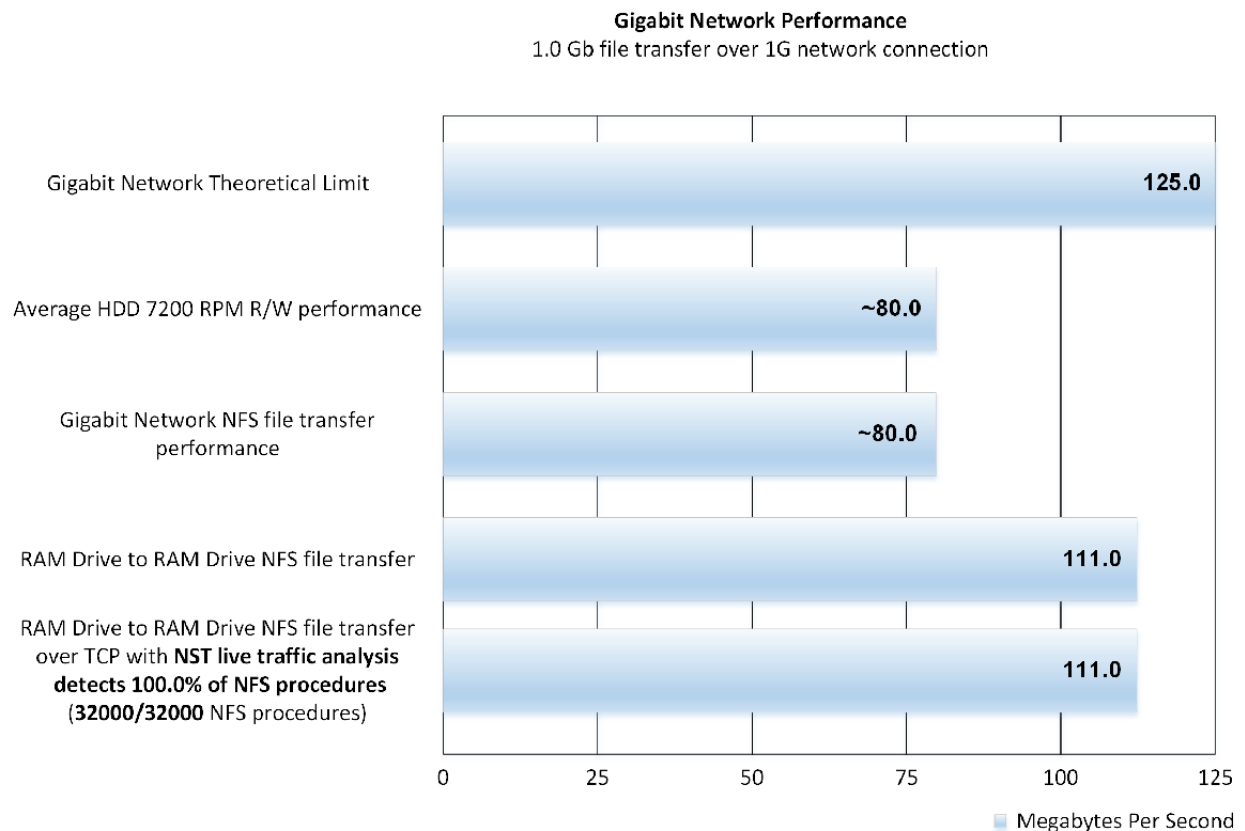
packets received by filtration: 2659749

packets dropped by kernel : 0

packets dropped by interface : 0

./nfstrace: Interrupt

We got 512913 WRITE procedures and 49 additional procedures on 76.9 Mb/s data transfer speed.



Picture 3 – nfstrace performance on Gigabit Ethernet

On Picture 3 you can see filtration performance of nfstrace tool during write 1G data from ram to ram drive over 1Gbps Ethernet.